



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich



## **Master's Thesis Nr. 164**

Systems Group, Department of Computer Science, ETH Zurich

A High-level Graph Query Language Interface for Differential Dataflow

by

Lukas Striebel

Supervised by

Prof. Timothy Roscoe, John Liagouris, Desislava Demitrova, Moritz Hoffmann

06.10.2016 - 06.04.2017



## **Abstract**

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Motivation . . . . .	3
<b>2</b>	<b>Background</b>	<b>4</b>
2.1	The Property Graph Query Language . . . . .	4
2.2	Differential Dataflow . . . . .	6
<b>3</b>	<b>Parser</b>	<b>7</b>
3.1	Path Clause . . . . .	8
3.2	Select Clause . . . . .	9
3.3	Where Clause . . . . .	10
3.3.1	Path Constraints . . . . .	10
3.3.2	Value Constraints . . . . .	10
3.4	Solution Modifier Clause . . . . .	11
<b>4</b>	<b>Query Evaluation</b>	<b>12</b>
4.1	Graph Loader . . . . .	13
<b>5</b>	<b>Benchmarking</b>	<b>14</b>
<b>6</b>	<b>Results and Discussion</b>	<b>15</b>
<b>7</b>	<b>Related Work</b>	<b>16</b>
7.1	PQL . . . . .	16
7.2	Green-Marl . . . . .	17
7.3	Gremlin . . . . .	18
7.4	SQLGraph . . . . .	19
7.5	GraphiQL . . . . .	20
<b>8</b>	<b>Summary</b>	<b>21</b>
8.1	Conclusion . . . . .	21
8.2	Acknowledgements . . . . .	21

# **1 Introduction**

## **1.1 Motivation**

## 2 Background

### 2.1 The Property Graph Query Language

The Property Graph Query Language (PGQL) was developed by Oskar van Rest at Oracle. [8] PGQL enables developers to write intuitive path queries over a property graph. A PGQL query consists of 4 clauses, two of them are optional.

- Path Clause (optional)
- Select Clause (required)
- Where Clause (required)
- Solution Modifier Clause (optional)

In the Path Clause custom path patterns are defined, which are to be used again in the Where Clause.

**Example:**

---

```
PATH connects_to := (:Generator) -[:has_connector]-> (:Connector WITH status =  
  'OPERATIVE') <-[:has_connector]- (:Generator)
```

---

The Select Clause bears great similarity to the SQL one. Here, the set of properties one wishes to retrieve is defined. Possible Selections are either

- everything, indicated by the use of \*
- certain attributes of edges and/or vertices
- Aggregation of certain attributes of edges and/or vertices.

Currently, there are five types of aggregations supported:

- COUNT, returns the number of tuples in the solution
- MAX, returns the maximum value of an attribute in any tuple. The specified attribute has to be numeric.
- MIN, returns the minimum value of an attribute in any tuple. The specified attribute has to be numeric.
- SUM, returns the sum of an attribute over all the tuples. The specified attribute has to be numeric.
- AVG, returns the average of an attribute over all the tuple. The specified attribute has to be numeric.

**Example:**

---

```
SELECT v.name, AVG(v.age)
```

---

The most complex part of the query is the Where Clause. In this section, all the requirements that the edges and vertices of the result set have to fulfill are specified.

## 2.2 Differential Dataflow

Differential Dataflow was developed by Frank McSherry at Microsoft. [5] Rather than rewriting the entire data every time a change occurs, Differential Dataflow only keeps track of the changes made to the data. This allows for huge timesavings when executing the same operation multiple times on different data, since the results of previous computations can be reused.



### 3 Parser

The Parser is build using nom, [1]. Nom was developed by Geoffroy Couprie. It is a byte oriented, zero copy, streaming Parser Library written in Rust. The library provides macros, functions and enums which facilitate the parsing process.

The most commonly used macros were:

- `do_parse!`: Takes a list of parsers as inputs and applies them sequentially, finally returns a tuple.
- `alt_complete!`: Takes a set of parsers as input and applies them until one succeeds. Returns the result of the first successful parser.
- `opt!`: Takes a parsers as input and makes it optional. Returns `None` if parser was unsuccessful or `Some` otherwise.
- `tag!`: Parses a specific String. Aborts if the String is not found.
- `named!`: Facilitates the creation of new custom parsers.
- `many0!`: Applies the parser 0 or more times.
- `many1!`: Applies the parser 1 or more times.

As mentioned in Section 2.1, the Parser has to recognize the following 4 clauses: Path definitions, select, where, solution modifier.

The main parse function therefore looks:

---

```
named!(pgql_query<Query>,  
  do_parse!(  
    paths: opt!(paths) >>  
    space >>  
    select: select_clause >>  
      space >>  
      vwhere: where_clause >>  
    space >>  
    solmod: opt!(solutionModifier) >>  
    (Query { select: select, vwhere: vwhere, paths: paths, solmod: solmod}))
```

---

### 3.1 Path Clause

The Path Clause is a list of definitions. Each Path definition starts with a name, followed by ‘:=’ and then a path description. The path defined in this clause can then be reused in the Where Clause. One of many difficulties encountered while writing the parser, is the ability to differentiate between a single path definition and multiple ones, separated by commas.

### 3.2 Select Clause

The Select Clause of PGQL is very similar to the SQL one. Started by the keyword ‘Select’, a list of attributes is provided. Attributes may be renamed with the keyword ‘as’. Aggregate functions like ‘Sum’, ‘Avg’ etc. may also be accessed.

### 3.3 Where Clause

The Where Clause consists of a list of Constraints. Each constraint defines either a path or value requirement that has to be fulfilled by the respective vertex or edge.

#### 3.3.1 Path Constraints

A Path Constraint requires a vertex to have certain edges to other vertices. For example, the constraint:  $(v) \rightarrow (u)$  requires that the vertex  $v$  has a direct edge to the vertex  $u$ .

#### 3.3.2 Value Constraints

Value Constraints are constraints on attributes of the Vertex, e.g.  $\text{name} = \text{'Alice'}$  or  $\text{age} < 40$ . Every Value Constraint has to include one or multiple vertex attributes, and one or more Literals. Literals are raw values, and come in 3 types:

- Strings e.g. 'Alice'
- Floats e.g. 40
- Booleans, which are either true or false

---

```
pub enum Literal {  
  Str(String),  
  Float(f32),  
  Boolean(bool),  
}  
  
named!(literal<Literal>,  
  alt_complete!(  
    float    => { |f| Literal::Float(f)          } |  
    boolean  => { |b| Literal::Boolean(b)         } |  
    string   => { |s| Literal::Str(String::from(s)) }  
  )  
);
```

---

### 3.4 Solution Modifier Clause

The Solution Modifier Clause consists of three parts:

- GroupBy Clause
- OrderClause
- Offset and Limit Clause

All clauses are optional and can be omitted.

## 4 Query Evaluation

## 4.1 Graph Loader

The Loading of the graph plays a major part of the entire program execution time. To load a graph from a text file into differential Dataflow, an entirely new parser had to be written.

## 5 Benchmarking



## 6 Results and Discussion

## 7 Related Work

### 7.1 PQL

A program query language, PQL for short, is a source language-independent notation to specify program queries and program views. PQL is used as an interface to Static Program Analyzers (SPA), interactive tools that enhance program understanding by answering queries about programs. Queris on global program design as well as searches for detail code patterns are both possible in PQL. Program queries and patterns supported by other notations described in literature and those supported by commercial tools can be written simply and naturally in PQL.[3]

## 7.2 Green-Marl

Green-Marl is a domain-specific language (DSL) with high level language construct that allow developers to describe their graph analysis algorithms intuitively, but expose the data-level parallelism inherent in the algorithms. Green-Marl comes with its own compiler which translates high-level algorithmic description written in Green-Marl into an efficient C++ implementation by exploiting this exposed datalevel parallelism. Furthermore, the Green-Marl compiler applies a set of optimizations that take advantage of the high-level semantic knowledge encoded in the Green-Marl DSL. Most graph analysis algorithms can be written very intuitively with Green-Marl and experimental results show that the compiler-generated implementation out of such descriptions performs just as well as or better than highly-tuned handcoded implementations.[2]

### 7.3 Gremlin

Developed by the Apache Software Foundation, Gremlin is a query language as well as a graph traversal machine. The graph traversal machine Gremlin consists of three parts that continuously interact with each other: first the graph, second the traversal and finally the set of traversers. The traversers move about the graph according to the instructions specified in the traversal, where the result of the computation is the ultimate locations of all halted traversers. A Gremlin machine can be executed over any supporting graph computing system such as an OLTP graph database and/or an OLAP graph processor. The language Gremlin is a functional language implemented in the user's native programming language. Gremlin supports both imperative and declarative querying. [6]

## 7.4 SQLGraph

SQLGraph is a Graph Store that combines existing relational optimizers with a novel schema, in an attempt to give better performance for property graph storage and retrieval than popular noSQL graph stores. The schema combines relational storage for adjacency information with JSON storage for vertex and edge attributes. This particular schema design has benefits compared to a purely relational or purely JSON solution. The query translation mechanism translates Gremlin queries with no side effects into SQL queries so that one can leverage relational query optimizers. [7]

## 7.5 GraphiQL

GRAPHiQL is an intuitive query language for graph analytics, which allows developers to reason in terms of nodes and edges rather than the tables and joins which are used in relational databases. GRAPHiQL provides key graph constructs such as looping, recursion, and neighborhood operations. At runtime, GRAPHiQL compiles graph programs into efficient SQL queries that can run on any relational database. [4]

## **8 Summary**

### **8.1 Conclusion**

### **8.2 Acknowledgements**

I would like to thank my supervisors John Liagouris, Desislava Dimitrova and Moritz Hoffmann.

## References

- [1] G. Couprie. Nom, a byte oriented, streaming, zero copy, parser combinators library in rust. In *2015 IEEE Security and Privacy Workshops*, pages 142–148, May 2015.
- [2] Sungpack Hong, Hassan Chafi, Edic Sedlar, and Kunle Olukotun. Green-marl: A dsl for easy and efficient graph analysis. In *Proceedings of the Seventeenth International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS XVII, pages 349–362, New York, NY, USA, 2012. ACM.
- [3] Stan Jarzabek. *PQL: A language for specifying abstract program views*, pages 324–342. Springer Berlin Heidelberg, Berlin, Heidelberg, 1995.
- [4] A. Jindal and S. Madden. Graphiq: A graph intuitive query language for relational databases. In *2014 IEEE International Conference on Big Data (Big Data)*, pages 441–450, Oct 2014.
- [5] Frank McSherry, Derek Murray, Rebecca Isaacs, and Michael Isard. Differential dataflow. In *Proceedings of CIDR 2013*, January 2013.
- [6] Marko A. Rodriguez. The gremlin graph traversal machine and language. *CoRR*, abs/1508.03843, 2015.
- [7] Wen Sun, Achille Fokoue, Kavitha Srinivas, Anastasios Kementsietsidis, Gang Hu, and Guotong Xie. Sqlgraph: An efficient relational-based property graph store. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, SIGMOD ’15, pages 1887–1901, New York, NY, USA, 2015. ACM.
- [8] Oskar van Rest, Sungpack Hong, Jinha Kim, Xuming Meng, and Hassan Chafi. Pgql: A property graph query language. In *Proceedings of the Fourth International Workshop on Graph Data Management Experiences and Systems*, GRADES ’16, pages 7:1–7:6, New York, NY, USA, 2016. ACM.